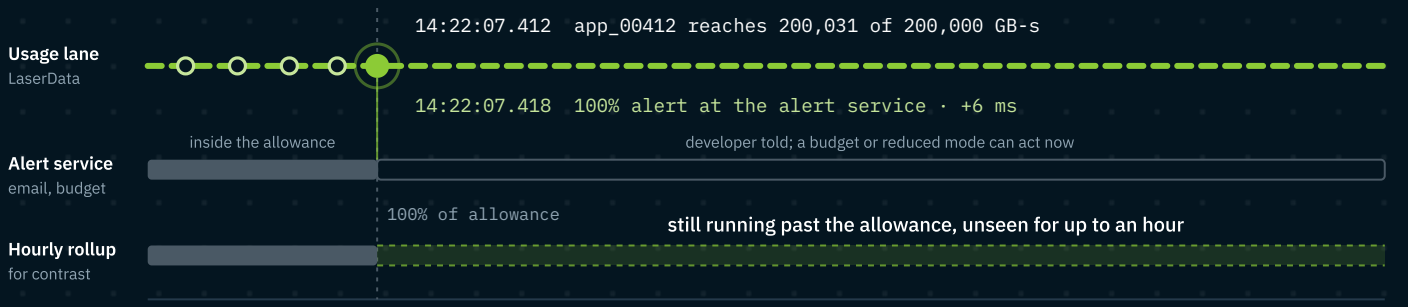
laserdata

Prepared for **Atlassian**

Real-time Forge usage alerts for Atlassian, on LaserData.

Every Forge invocation's usage reaches **its app's allowance check, and every crossing reaches the alert service, in milliseconds**, counted once. It runs beside StreamHub, so nothing else has to move.

One Forge app at its allowance,
as the event log sees it.
Now multiply by every app.




```
14:22:07.398   usage.total     app_00412  rt-03  app at 199,994 GB-s
14:22:07.412   usage.total     app_00412  rt-06  200,031 GB-s · rt-03 retry adds 0
14:22:07.416   alert.crossing  app_00412  100% of allowance  +4 ms
14:22:07.418   alert service   email to contributors, budget check  +2 ms
```

From the creators of APACHE IGGY

15×
faster P99 latency
than Apache Kafka

4×
higher throughput
than Apache Kafka

60–70%
lower infrastructure cost
compute, storage and network

Apache Iggy vs Apache Kafka, LaserData testing. Results vary by workload.

What you shared	Forge usage runs **~300M events a day, heading to 1B**. The rules: no double counting, a usageUniqueId so retries add nothing, alerts at 50/75/90/100%. Developers ask for budgets the platform "doesn't support yet". (Forge billing post, Apr 2026; Forge community, Aug 2026)
Where it fits	**The usage lane:** Forge usage keyed by app into the allowance tracker, crossings out to the alert service. StreamHub, UTS, Plato and Commerce stay.
Next	**AI agents** act on the same lane, with session, context, memory and replay from Laser SDK.
Underneath	**One durable event log.** LaserData, powered by Apache Iggy.

Try it in 15 minutes. Let the numbers decide.

A ready-made kit for Atlassian: **8 runtime hosts, 5,000 apps, 4,000 invocations/s** (~350M usage events a day), an allowance tracker with the four Forge levels and a budget level, retries sent twice, a busy hour, a product-analytics flood next door and a StreamHub-subscription shadow bridge. It prints your lane's p99, alerts per level and retries absorbed every second.

Open the Atlassian kit and write-up →

Then a 2–4 week evaluation on one mirrored usage subscription, alongside StreamHub. **No advantage, no migration.**

kranti@laserdata.com laserdata.com

© 2026 LaserData, Inc. Apache, Apache Iggy and the Iggy logo are trademarks of The Apache Software Foundation; their use does not imply endorsement by the ASF. The Atlassian name identifies the intended reader only; no affiliation is implied. Facts are from Atlassian's engineering posts "Engineering the Forge billing platform for reliability and scale" (Apr 2026) and "Scaling StreamHub" (2026), its Forge pricing and usage-alerts docs (read Oct 2026) and an Atlassian reply in the Forge developer community (Aug 2026). The app, host names, timestamps and the hourly rollup are illustrative.