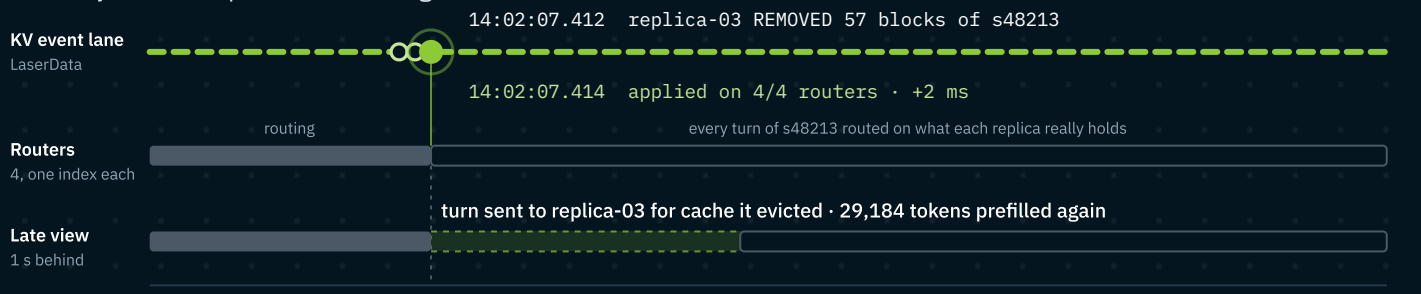
laserdata

Prepared for **Baseten**

Real-time KV cache events for Baseten, on LaserData.

Every block a replica stores or evicts reaches **every KV cache-aware router in milliseconds**, before the next request on that prefix. It runs beside your event plane, so nothing else has to move.

One evicted conversation,
as the event log sees it.
Now multiply by every replica.




```
14:02:07.412   kv.removed   replica-03  57 blocks of s48213, under memory pressure
14:02:07.414   kv.applied   4/4 routers  replica-03 now holds 41 of its blocks  +2 ms
14:02:07.430   request      s48213 turn 12, 50,176 tokens  → replica-05, 72 cached
1 s behind     → replica-03 for 98 cached blocks, 41 real: 29,184 tokens prefilled again
```

From the creators of APACHE IGGY

15×
faster P99 latency
than Apache Kafka

4×
higher throughput
than Apache Kafka

60–70%
lower infrastructure cost
compute, storage and network

Apache Iggy vs Apache Kafka, LaserData testing. Results vary by workload.

What you shared	Baseten is a heavy user of Dynamo's KV cache-aware routing and **"often see[s] 2x faster TTFT in production"**; on shadowed traffic P95 and P99 fell 48–49%. The Inference Platform team owns routing, autoscaling, metering and quotas. (Baseten blog, Oct 2025; job post, Oct 2026)
Where it fits	**The KV event lane:** every block stored or evicted, keyed by deployment, to every router, in order and without gaps; a new router rebuilds from the log. Dynamo, your engines and routing logic stay.
Next	**Usage and telemetry** for metering and quotas on the same log, with session, context and replay from Laser SDK.
Underneath	**One durable event log.** LaserData, powered by Apache Iggy.

Try it in 15 minutes. Let the numbers decide.

A ready-made kit for Baseten: **16 replicas, 4 KV cache-aware routers, 6,000 long-context sessions at 600 requests/s**, a router that joins mid-run and rebuilds from the log, a telemetry flood next door and a ZMQ shadow bridge. It prints your lane's p99, stale routes and tokens prefilled again every second.

Open the Baseten kit and write-up →

Then a 2–4 week evaluation on one deployment's KV events, mirrored beside your event plane. **No advantage, no migration.**

kranti@laserdata.com laserdata.com

© 2026 LaserData, Inc. Apache, Apache Iggy and the Iggy logo are trademarks of The Apache Software Foundation; their use does not imply endorsement by the ASF. The Baseten name identifies the intended reader only; no affiliation is implied. Facts are from Baseten's blog posts "2x faster inference with KV cache-aware routing" (Oct 2025, updated Mar 2026) and "The Baseten Inference Stack at NVIDIA Dynamo Day" (Feb 2026), its job post "Software Engineer - Inference Platform" (Oct 2026) and NVIDIA Dynamo's router documentation (read Oct 2026). Session, replicas, timestamps, block counts and the 1-second lag are illustrative.