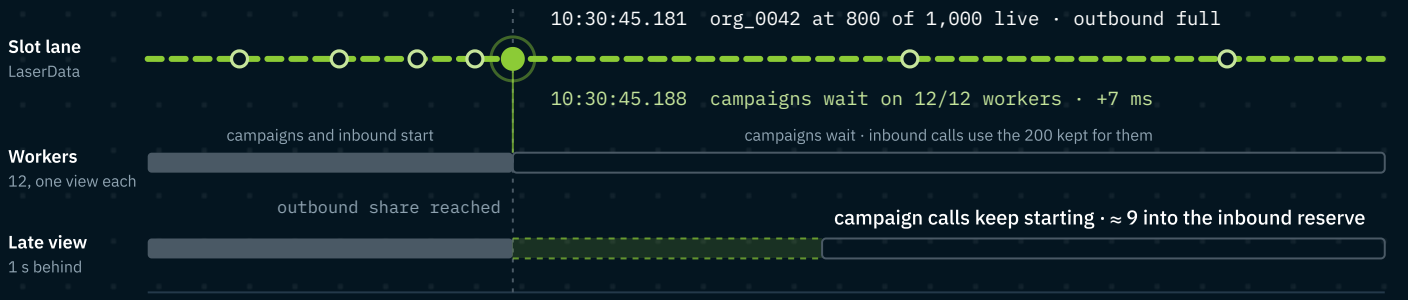
laserdata

Prepared for **Sarvam AI**

Real-time call concurrency for Sarvam AI, on LaserData.

Every conversation start and end reaches **the plan-ceiling ledger, and every headroom change reaches every worker, in milliseconds**, even at peak. It runs beside your call stack, so nothing else has to move.

One organization at its ceiling,
as the event log sees it.
Now multiply by every org running a campaign.




```
10:30:45.061   slot.live   org_0042  wkr-07  end        799/1,000
10:30:45.181   slot.live   org_0042  wkr-03  campaign   800/1,000
10:30:45.185   headroom    org_0042  OUTBOUND_FULL  +4 ms
10:30:45.188   applied     12/12 workers, before the next conversation  +3 ms
```

From the creators of APACHE IGGY

15×
faster P99 latency
than Apache Kafka

4×
higher throughput
than Apache Kafka

60–70%
lower infrastructure cost
compute, storage and network

Apache Iggy vs Apache Kafka, LaserData testing. Results vary by workload.

What you shared	Plan ceilings on live conversations are **"shared across everything that dials"**: campaigns, instant outbound and inbound. Past one, "new calls are queued or rejected". Millions of voice calls a day, at sub-second latency SLOs. (Sarvam docs, 2026; job posts, 2026)
Where it fits	**The concurrency lane:** conversation starts and ends keyed by organization into the plan-ceiling ledger, headroom back to every worker. Telephony, Temporal, PostgreSQL and Redis stay.
Next	**AI agents** act on the same lane, with session, context, memory and replay from Laser SDK.
Underneath	**One durable event log.** LaserData, powered by Apache Iggy.

Try it in 15 minutes. Let the numbers decide.

A ready-made kit for Sarvam AI: **12 voice-agent workers, 300 organizations, 60 conversations/s** at 90 s talk time, a ledger with an inbound reserve, a campaign-launch spike, a post-call payload flood next door and a shadow bridge for your call events. It prints your lane's p99 and calls started past a ceiling every second.

Open the Sarvam AI kit and write-up →

Then a 2–4 week evaluation on one mirrored stream of call events, alongside your current stack. **No advantage, no migration.**

kranti@laserdata.com laserdata.com

© 2026 LaserData, Inc. Apache, Apache Iggy and the Iggy logo are trademarks of The Apache Software Foundation; their use does not imply endorsement by the ASF. The Sarvam AI name identifies the intended reader only; no affiliation is implied. Facts are from Sarvam's public Voice Agents docs ("Dialing rate & concurrency" and "Campaigns", read Oct 2026) and its job posts "Staff Engineer - Telephony" and "Staff Engineer, Product Infrastructure" (Oct 2026). Organization, ceiling, inbound reserve, workers, timestamps and the one-second lag are illustrative.